

Göttingen



Timeline

Lecture
3.3. - 7.3.

8-12h

Your Talks:
planned projects
7.3.

Lab Course
10.3. - 28.3.

24/7

Talks/Posters:
your results
??4.

Goals

- learn CUDA
- learn to use CUDA-based libraries
- use CUDA in a C++ environment
- convert project into article (cf. step-16)

Modules

B.Mat **106**

Grundlagen Wiss. Rechnen

B.Mat.**730**

Weiterführung Praktikum
Wiss. Rechnen

BSc.
Computer
Science

BSc.
Math,
Physics

Masters: ?

Projects

Each participant works on a larger project which he/she presents before the lab course starts. The talk should be about 15-20 minutes and cover:

- what will be simulated
 - which part of the problem is supposed to profit from parallelization
 - sketch of algorithm
 - algorithmic problems due to parallelization (some algorithms simply cannot be parallelized)
-
- quantum time propagation
 - matrix assembly for integral equations (FEM-BEM coupling)
 - super-resolution algorithms for optical microscopy
 - X-ray physics (all sorts of phase reconstruction)

Example: 2013's Lecture

	Mo	Di	Mi	Do	Fr
8:15 - 9:45	CUDA Basics programming model, hardware, memory transfers, C++	CUDA Advanced more memory handling, OpenGL, PAAL	Fluid Dynamics <i>A. Tilgner</i>	Finite Differences Multigrid theory, implementation	Rashmi Barbate Prefix Sums
					Sanjeev Laha Histograms
					Simon Maretzke GMT, Bessel, ...
					break GPU-assisted routing
10:00 - 11:30	Linear Algebra sparse matrices, deal.II interoperability	C++ Expression Templates SciPal, dense LA	Finite Elements C(ontinuous) <i>G. Lube</i>	Boundary Elements algorithms and implementation	Moritz Doll Searching/Sorting
					Simon Schütz R, ...
					break Subasanth Gosh Histograms
					Pranay Tare Preconditioning / Multigrid
12:00 - 14:30	Project Structure svn, documenting, QTCreator, coding conventions	Debugging Tools for analysing CUDA code	QThreads		Christian Holme BEM Assembly
					break
		OpenMP			

Lab Course

Structure of programs:

This will be the essence
of your introductory talk



Introduction
(2-3 pages A4, if printed)

Commented
Program

This will be the **basis**
for your final talk/poster



Results
colorful pictures

- Core time: 9-16 h
- Show up each day for discussions, etc ...
- Stick to the prepared structure of the steps

Further details:

<http://num.math.uni-goettingen.de/~stkramer/doc/tutorial/index.html>

Presentation of Results

- to be scheduled for late april
- either seminar or poster session, depends on the number of participants

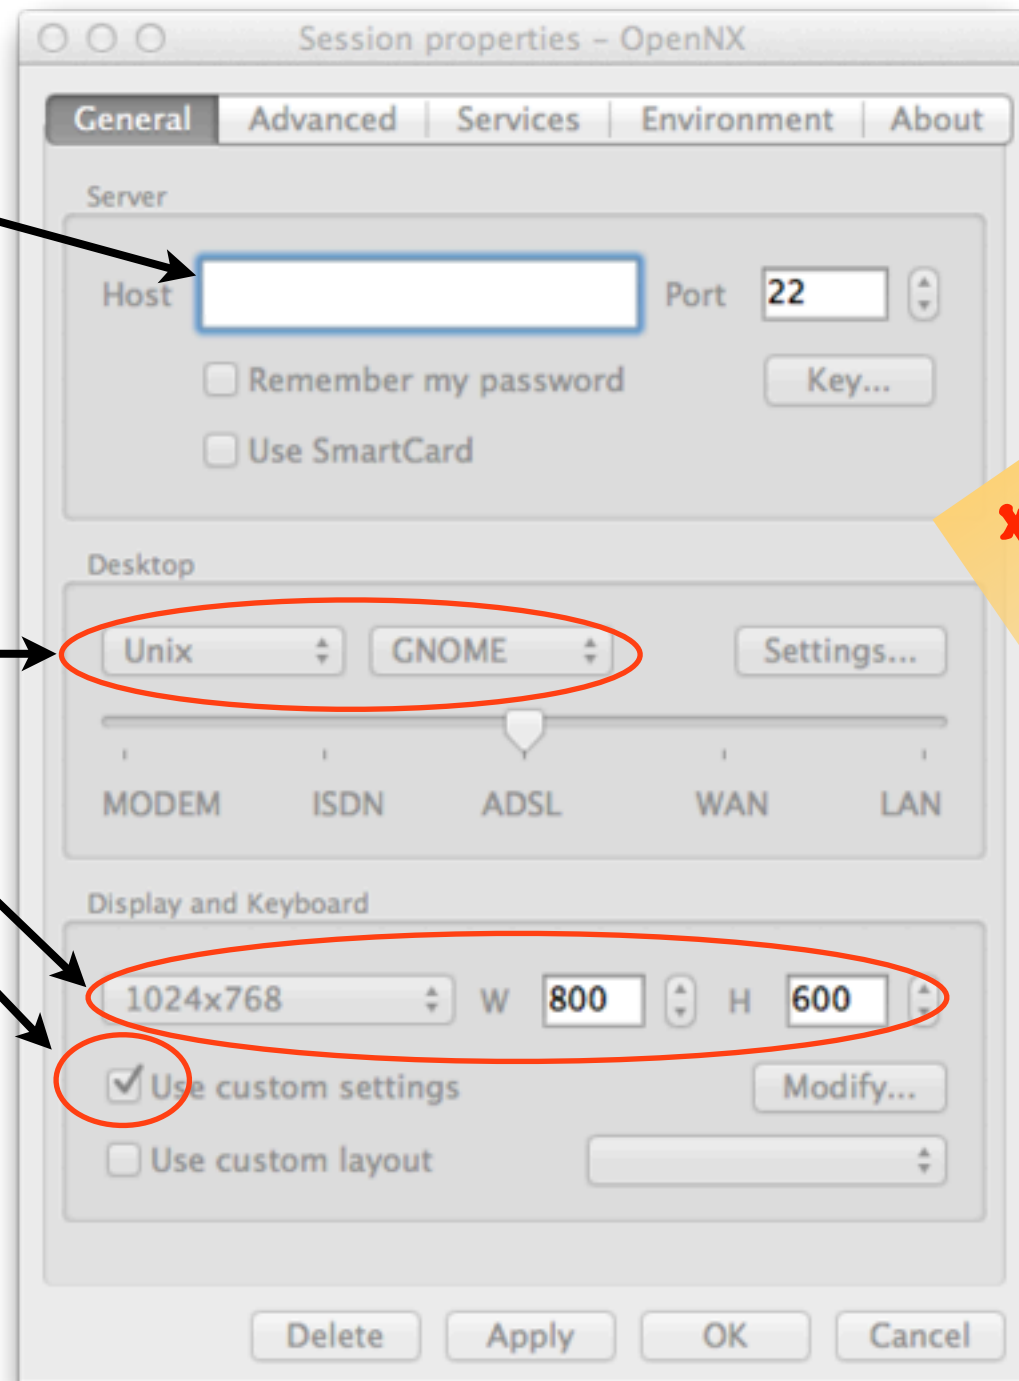
Lab Course

How To

Login via NXClient/OpenNX

enter the IP I
emailed you

window manager
config as shown
here

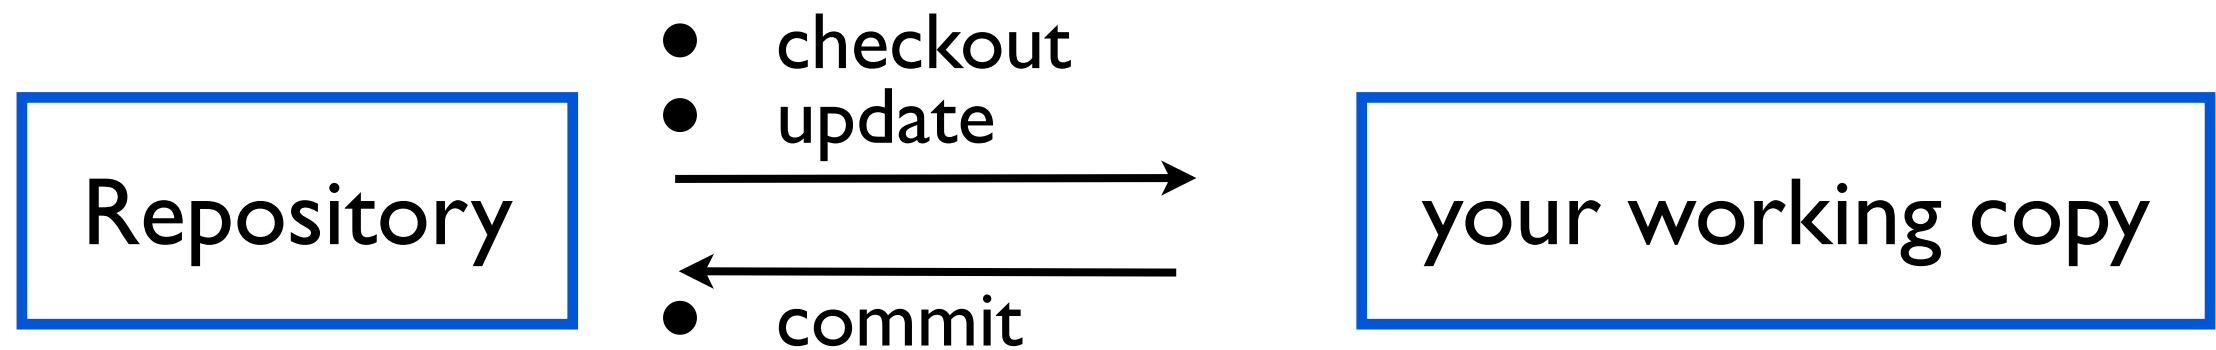


this still works, but x2go is
the preferred way. Works
essentially the same way

Then:
log in



svn



repository location (once it exists & you got access):

https://svn.num.math.uni-goettingen.de/svn/cuda/Praktikum_2014

location of working copy: *somewhere in your home directory*

how to create it:

1. open a terminal
2. create a folder 'cuda' in your home directory:

```
mkdir ~/cuda
```

3. go to that directory:

```
cd ~/cuda
```

4. check out the latest version of the whole lab course:

```
svn checkout https://svn.num.math.uni-goettingen.de/svn/cuda/Praktikum\_2014
```

5. once you have changed something which you want to keep, commit it to the repository:

```
cd ~/cuda/Praktikum_2014/<some subdir>
```

```
svn commit -m "<my commit message>"
```

At some point you have to
enter your account password
for the svn access

Your Project

location :

`Praktikum_2014/testsite/step-<xy>`

`<xy>` will be some number

central configuration file:

`step-<xy>/step-<xy>.pro`

Project Structure

host side

- compiled by g++

step-xy.cpp

main source
file

cuda_driver_step-xy.h*

header file for
driver class

- manages CUDA-based computations
- takes care of memory transfer

cuda_kernel_wrapper_step-xy.cu.h

header file for

interface template class

- contains all wrapper functions which call kernels
- can be used to abstract from parallelization technique

cuda.h, cuda_runtime_api.h, etc ... should go here

These should not
reference any CUDA
header like cuda.h,
cuda_runtime_api.h,
etc ...

device side

- compiled by nvcc

cuda_kernel_step-xy.cu

source file seen by nvcc

cuda_kernel_step-xy.cu.c

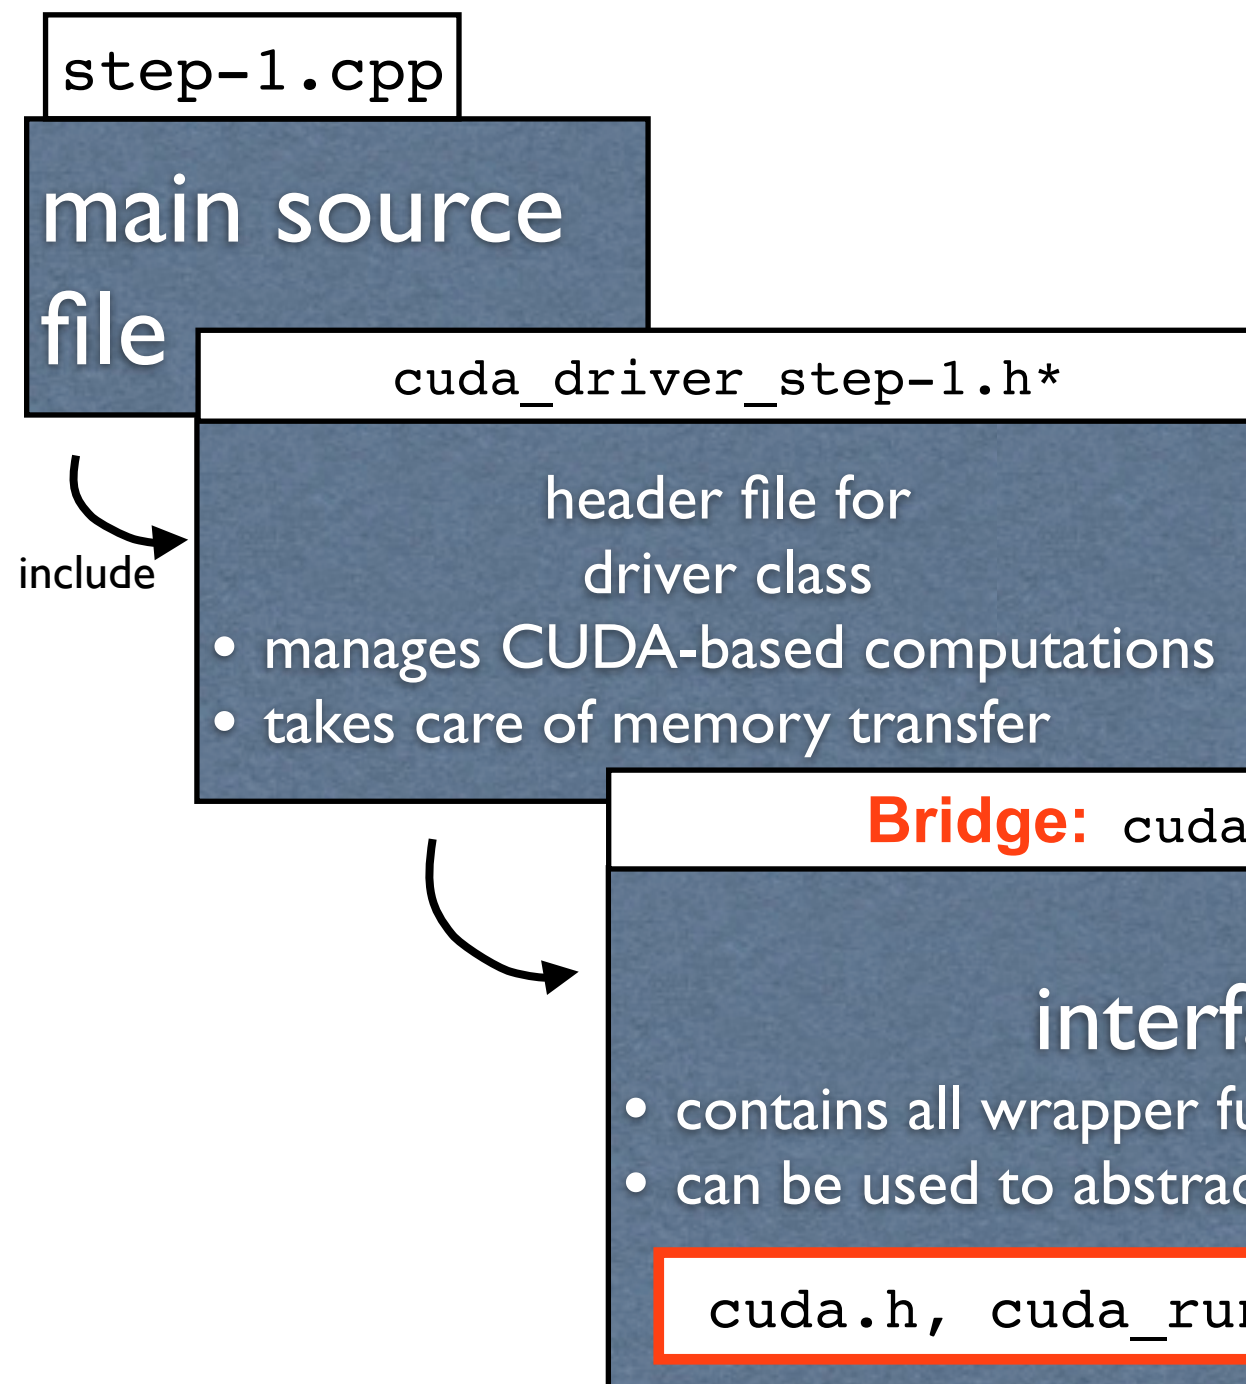
source file in disguise; contains:

- device functions
- device classes
- kernels
- template specializations of wrapper class

include

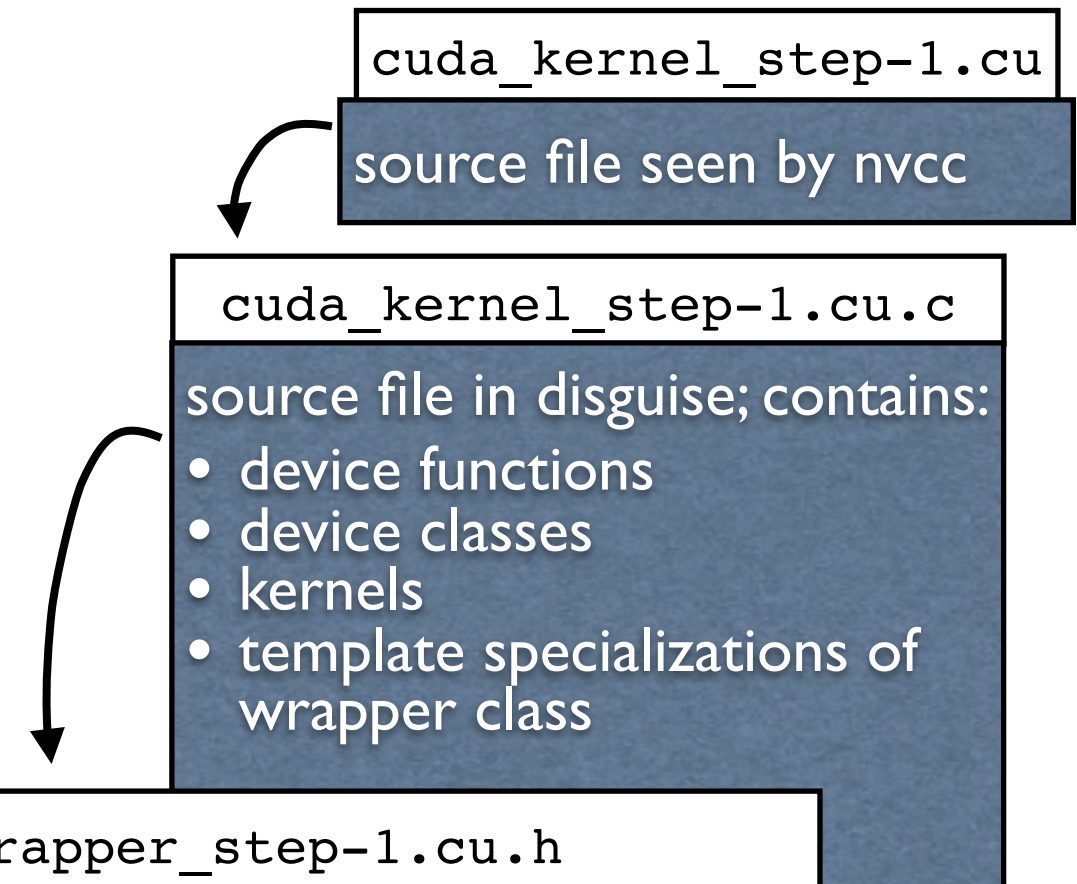
host side

- compiled by g++



device side

- compiled by nvcc



QtCreator

- can be found in Ubuntu's Application/Developer menu
- to work on your project, open `step-<xy>.pro`
- the “Shadow build” configuration can be found in the project settings of QtCreator

- **to avoid lengthy paths for the shadow build directories change the setting in**

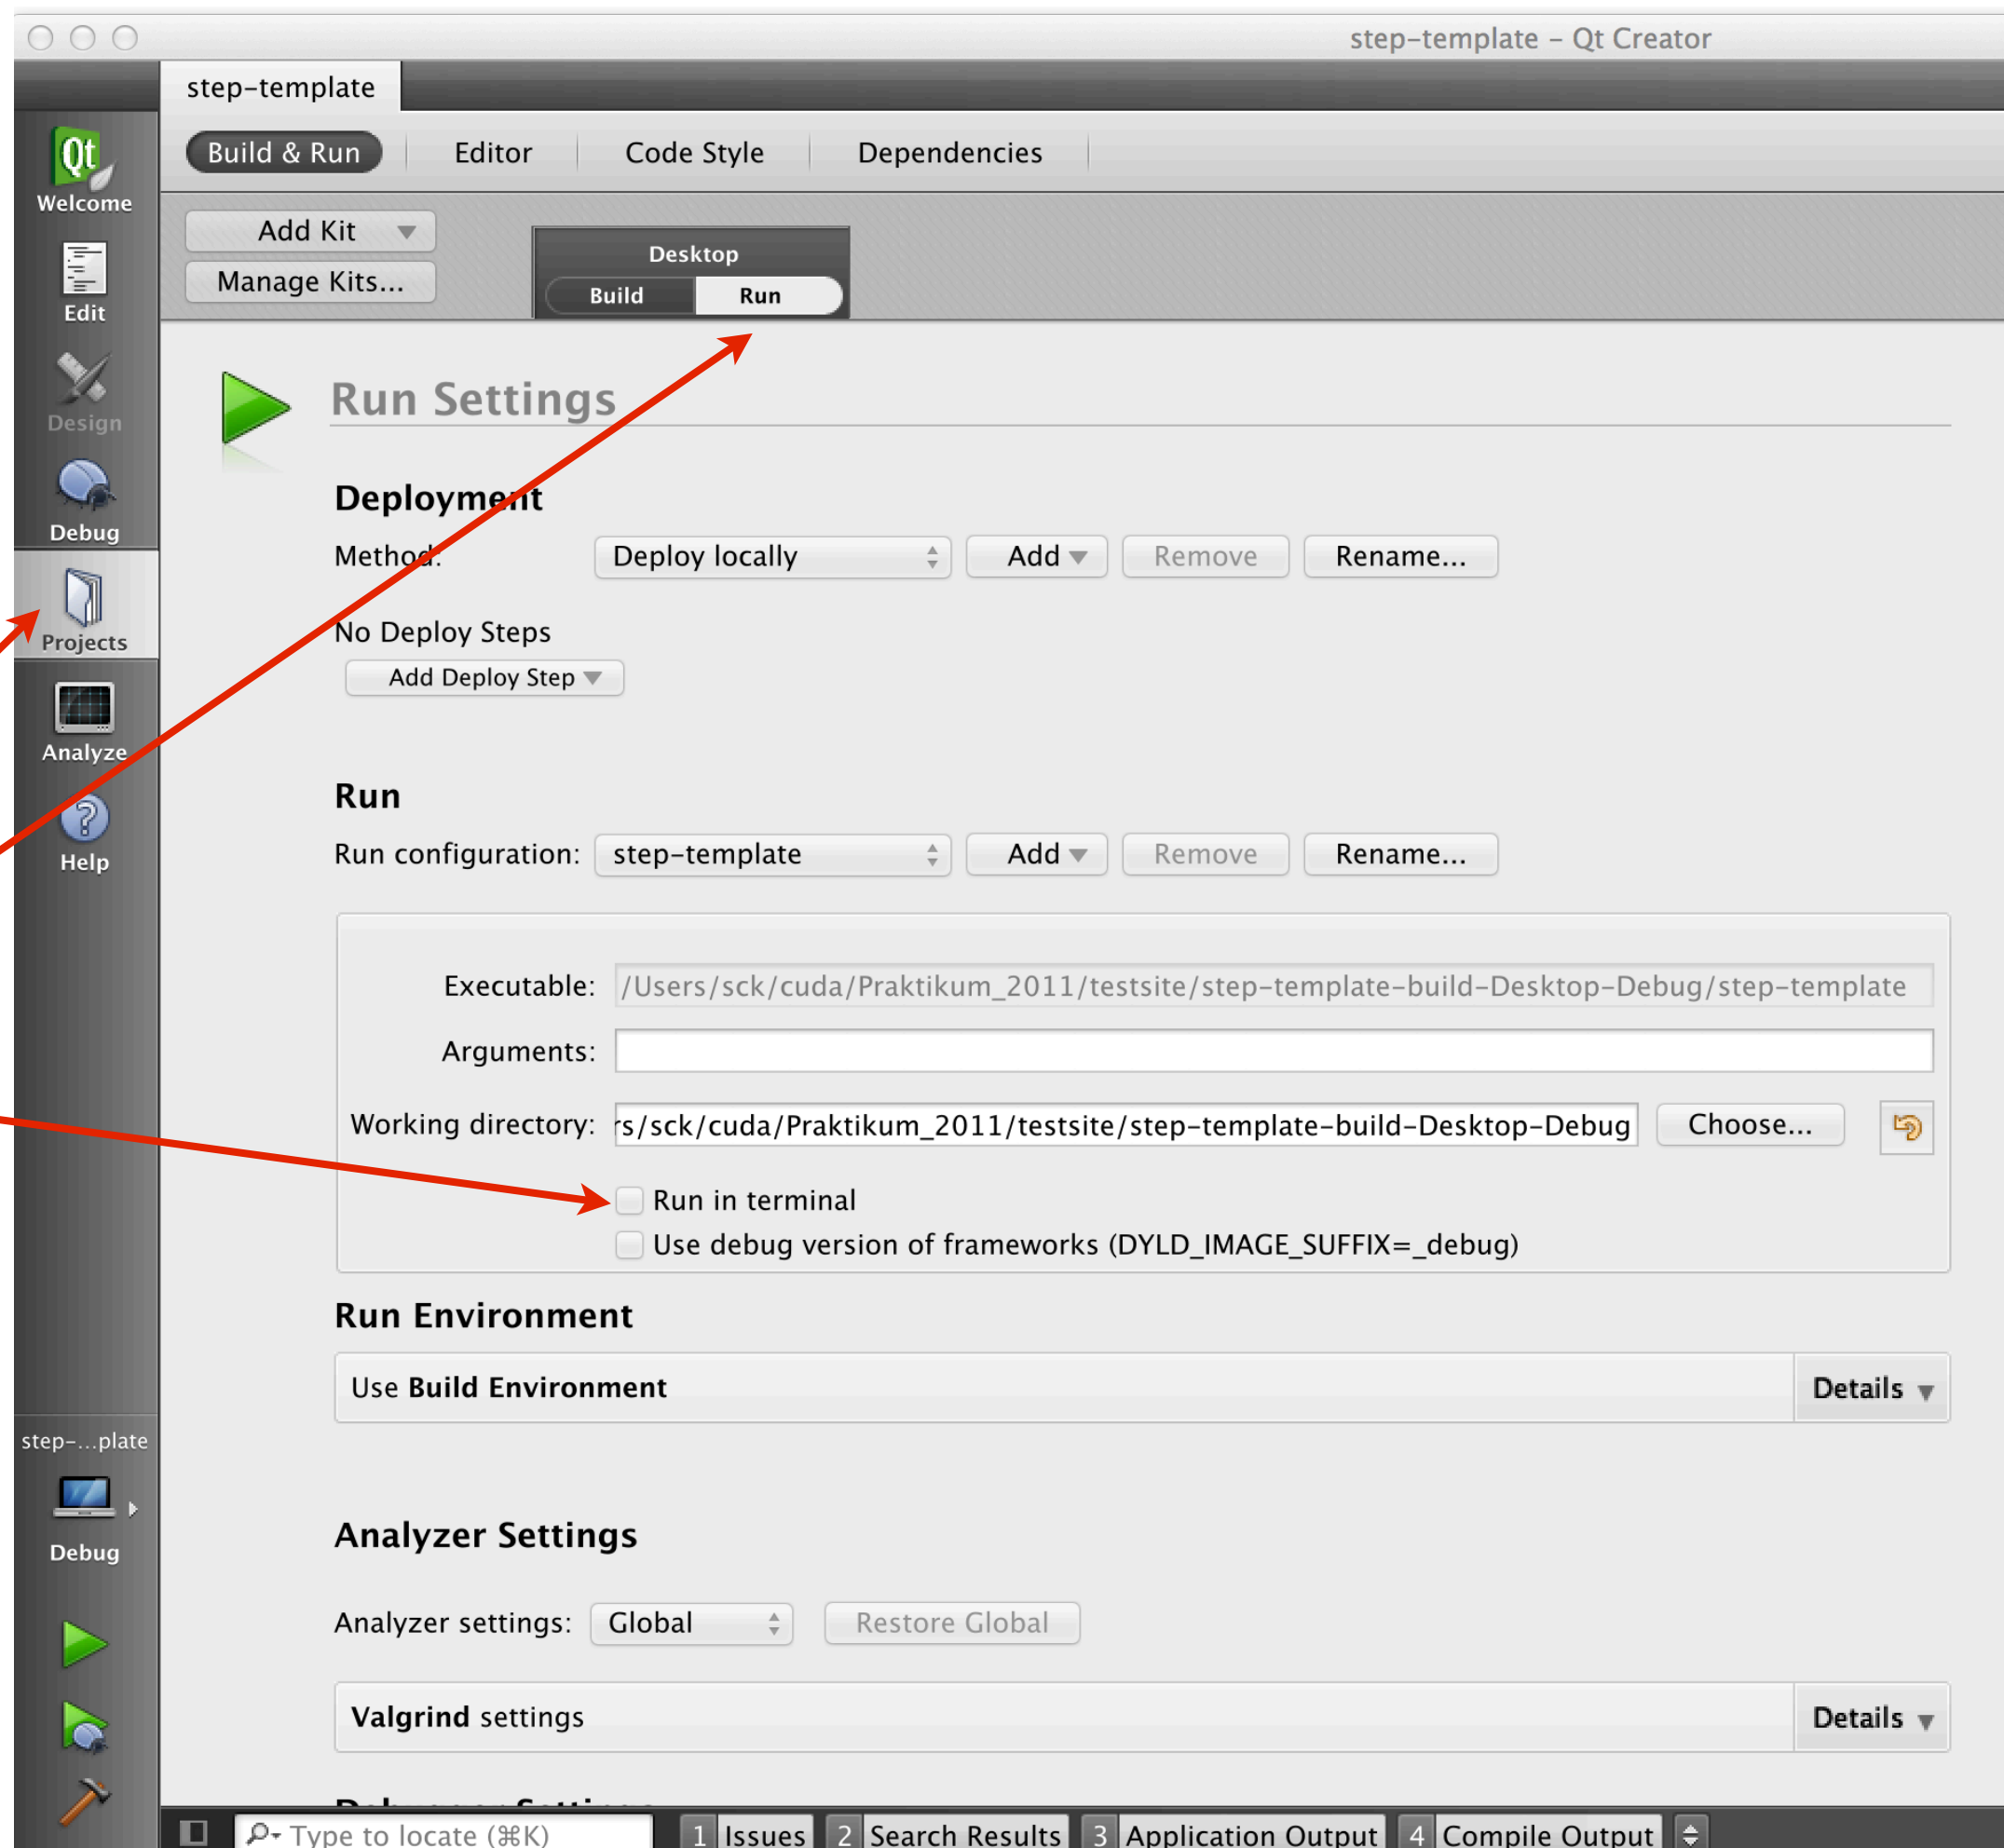
Preferences/Build&Run/Default Build Directory

- to run your program from within QtCreator, disable “run in terminal” (see next page)

1.

2.

3.



Generate Project's Doc

go to :

```
Praktikum_2014/scripts/doc
```

run script:

```
./make_step_doc.py <xy>
```

<xy> is the
number of your
project

open in a browser:

```
firefox ../../doc/tutorial/index.html
```

... and look for your number in the
navigation bar on the left

for further details:

`Praktikum_2014/readme.html`